

Machine learning - Assignment 1

Name: Yudhveer Singh Runnoo

Reg. no.: 22MCA0421

Q1. Design a simple linear regression model using all possible features and find the minimum MSE and R2 Score

```
In [1]: import matplotlib.pyplot as plt
import pandas as pd
import pylab as pl
import numpy as np
%matplotlib inline
from sklearn import linear_model
from sklearn.metrics import r2_score
```

```
In [2]: df = pd.read_csv("FuelConsumption.csv")
# take a look at the dataset
df.head()
```

```
Out[2]:
```

	MODELYEAR	MAKE	MODEL	VEHICLECLASS	ENGINE SIZE	CYLINDERS	TRANSMISSION	FUELTYPE
0	2014	ACURA	ILX	COMPACT	2.0	4		AS5
1	2014	ACURA	ILX	COMPACT	2.4	4		M6
2	2014	ACURA	ILX HYBRID	COMPACT	1.5	4		AV7
3	2014	ACURA	MDX 4WD	SUV - SMALL	3.5	6		AS6
4	2014	ACURA	RDX AWD	SUV - SMALL	3.5	6		AS6

```
In [3]: # summarize the data
df.describe()
```

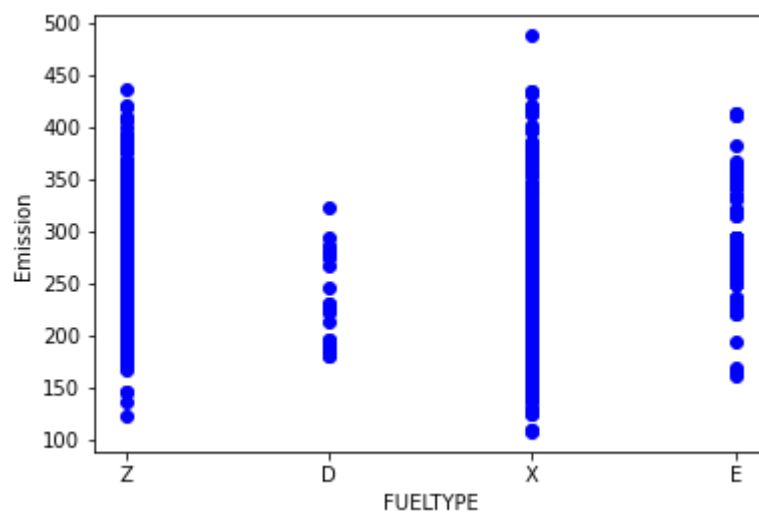
```
Out[3]:
```

	MODELYEAR	ENGINE SIZE	CYLINDERS	FUELCONSUMPTION_CITY	FUELCONSUMPTION_HWY
count	1067.0	1067.000000	1067.000000	1067.000000	1067.000000
mean	2014.0	3.346298	5.794752	13.296532	9.474602
std	0.0	1.415895	1.797447	4.101253	2.794510
min	2014.0	1.000000	3.000000	4.600000	4.900000
25%	2014.0	2.000000	4.000000	10.250000	7.500000
50%	2014.0	3.400000	6.000000	12.600000	8.800000
75%	2014.0	4.300000	8.000000	15.550000	10.850000
max	2014.0	8.400000	12.000000	30.200000	20.500000

```
In [4]: pd.unique(df['FUELTYPE'])
```

```
Out[4]: array(['Z', 'D', 'X', 'E'], dtype=object)
```

```
In [5]: plt.scatter(df.FUELTYPE, df.CO2EMISSIONS, color='blue')
plt.xlabel("FUELTYPE")
plt.ylabel("Emission")
plt.show()
```




```
In [6]: cdf = df[['ENGINE_SIZE', 'CYLINDERS', 'FUELCONSUMPTION_COMB', 'CO2EMISSIONS']]
cdf.head(9)
```

```
Out[6]:
```

	ENGINE_SIZE	CYLINDERS	FUELCONSUMPTION_COMB	CO2EMISSIONS
0	2.0	4	8.5	196
1	2.4	4	9.6	221
2	1.5	4	5.9	136
3	3.5	6	11.1	255
4	3.5	6	10.6	244
5	3.5	6	10.0	230
6	3.5	6	10.1	232
7	3.7	6	11.1	255
8	3.7	6	11.6	267

```
In [7]: plt.scatter(cdf.CYLINDERS, cdf.CO2EMISSIONS, color='blue')
plt.xlabel("CYLINDERS")
plt.ylabel("Emission")
plt.show()
```



```
In [8]: msk = np.random.rand(len(df)) < 0.8
train = cdf[msk]
test = cdf[~msk]
print(train)
print(test)
```

	ENGINE SIZE	CYLINDERS	FUEL CONSUMPTION_COMB	CO2 EMISSIONS
0	2.0	4	8.5	196
1	2.4	4	9.6	221
2	1.5	4	5.9	136
3	3.5	6	11.1	255
4	3.5	6	10.6	244
...
1061	3.2	6	11.2	258
1062	3.0	6	11.8	271
1064	3.0	6	11.8	271
1065	3.2	6	11.3	260
1066	3.2	6	12.8	294

[833 rows x 4 columns]

	ENGINE SIZE	CYLINDERS	FUEL CONSUMPTION_COMB	CO2 EMISSIONS
5	3.5	6	10.0	230
8	3.7	6	11.6	267
20	2.0	4	10.0	230
21	2.0	4	9.3	214
22	2.0	4	10.0	230
...
1041	2.0	4	6.9	186
1048	2.0	4	7.1	192
1054	3.6	6	12.2	281
1060	3.0	6	11.5	264
1063	3.2	6	11.5	264

[234 rows x 4 columns]

```
In [9]: regr = linear_model.LinearRegression()
train_x = np.asanyarray(train[['ENGINE SIZE']])
train_y = np.asanyarray(train[['CO2 EMISSIONS']])
regr.fit (train_x, train_y)
# The coefficients
print ('Coefficients: ', regr.coef_)
print ('Intercept: ',regr.intercept_)
```

```
Coefficients: [[38.93471163]]
Intercept: [125.52787705]
```

```
In [10]: test_x = np.asanyarray(test[['ENGINE SIZE']])
test_y = np.asanyarray(test[['CO2 EMISSIONS']])
#print(test_y)
```

```

test_y_ = regr.predict(test_x)
#print(test_y_)
print("Mean absolute error: %.2f" % np.mean(np.absolute(test_y_ - test_y)))
print("Residual sum of squares (MSE): %.2f" % np.mean((test_y_ - test_y) ** 2))
print("R2-score: %.2f" % r2_score(test_y_ , test_y) )

```

Mean absolute error: 24.39
Residual sum of squares (MSE): 1050.69
R2-score: 0.64

```

In [11]: regr = linear_model.LinearRegression()
train_x = np.asanyarray(train[['CYLINDERS']])
train_y = np.asanyarray(train[['CO2EMISSIONS']])
regr.fit (train_x, train_y)
# The coefficients
print ('Coefficients: ', regr.coef_)
print ('Intercept: ',regr.intercept_)

```

Coefficients: [[29.98316342]]
Intercept: [82.45502337]

```

In [12]: test_x = np.asanyarray(test[['CYLINDERS']])
test_y = np.asanyarray(test[['CO2EMISSIONS']])
#print(test_y)
test_y_ = regr.predict(test_x)
#print(test_y_)
print("Mean absolute error: %.2f" % np.mean(np.absolute(test_y_ - test_y)))
print("Residual sum of squares (MSE): %.2f" % np.mean((test_y_ - test_y) ** 2))
print("R2-score: %.2f" % r2_score(test_y_ , test_y) )

```

Mean absolute error: 26.04
Residual sum of squares (MSE): 1112.77
R2-score: 0.63

Q 2. Develop a multiple linear regression (MLR) using more than one feature and obtain the minimum possible error.

```
In [1]: import matplotlib.pyplot as plt
import pandas as pd
import pylab as pl
import numpy as np
%matplotlib inline
from sklearn import linear_model
from sklearn.metrics import r2_score
```

```
In [2]: df = pd.read_csv("FuelConsumption.csv")
# take a look at the dataset
df.head()
```

```
Out[2]:
```

	MODELYEAR	MAKE	MODEL	VEHICLECLASS	ENGINE SIZE	CYLINDERS	TRANSMISSION	FUELTYPE
0	2014	ACURA	ILX	COMPACT	2.0	4	AS5	
1	2014	ACURA	ILX	COMPACT	2.4	4	M6	
2	2014	ACURA	ILX HYBRID	COMPACT	1.5	4	AV7	
3	2014	ACURA	MDX 4WD	SUV - SMALL	3.5	6	AS6	
4	2014	ACURA	RDX AWD	SUV - SMALL	3.5	6	AS6	

```
In [3]: # summarize the data
df.describe()
```

```
Out[3]:
```

	MODELYEAR	ENGINE SIZE	CYLINDERS	FUELCONSUMPTION_CITY	FUELCONSUMPTION_HWY
count	1067.0	1067.000000	1067.000000	1067.000000	1067.000000
mean	2014.0	3.346298	5.794752	13.296532	9.474602
std	0.0	1.415895	1.797447	4.101253	2.794510
min	2014.0	1.000000	3.000000	4.600000	4.900000
25%	2014.0	2.000000	4.000000	10.250000	7.500000
50%	2014.0	3.400000	6.000000	12.600000	8.800000
75%	2014.0	4.300000	8.000000	15.550000	10.850000
max	2014.0	8.400000	12.000000	30.200000	20.500000


```
In [4]: cdf = df[['ENGINE SIZE', 'CYLINDERS', 'FUELCONSUMPTION_COMB', 'CO2EMISSIONS']]
cdf.head(9)
```

```
Out[4]:
```

	ENGINE SIZE	CYLINDERS	FUELCONSUMPTION_COMB	CO2EMISSIONS
0	2.0	4	8.5	196
1	2.4	4	9.6	221
2	1.5	4	5.9	136

	ENGINE SIZE	CYLINDERS	FUEL CONSUMPTION_COMB	CO2 EMISSIONS
3	3.5	6	11.1	255
4	3.5	6	10.6	244
5	3.5	6	10.0	230
6	3.5	6	10.1	232
7	3.7	6	11.1	255
8	3.7	6	11.6	267

```
In [5]: viz = cdf[['CYLINDERS', 'ENGINE SIZE', 'CO2 EMISSIONS', 'FUEL CONSUMPTION_COMB']]
viz.hist()
plt.show()
```



```
In [6]: msk = np.random.rand(len(df)) < 0.8
train = cdf[msk]
test = cdf[~msk]
print(train)
print(test)
```

	ENGINE SIZE	CYLINDERS	FUEL CONSUMPTION_COMB	CO2 EMISSIONS
0	2.0	4	8.5	196
1	2.4	4	9.6	221
2	1.5	4	5.9	136
3	3.5	6	11.1	255
4	3.5	6	10.6	244
...
1060	3.0	6	11.5	264
1061	3.2	6	11.2	258
1062	3.0	6	11.8	271
1064	3.0	6	11.8	271
1065	3.2	6	11.3	260

[850 rows x 4 columns]

	ENGINE SIZE	CYLINDERS	FUEL CONSUMPTION_COMB	CO2 EMISSIONS
5	3.5	6	10.0	230
8	3.7	6	11.6	267
12	5.9	12	15.6	359
13	5.9	12	15.6	359
40	2.0	4	9.2	212
...
1032	2.0	4	7.2	194
1042	1.4	4	5.4	124
1048	2.0	4	7.1	192
1063	3.2	6	11.5	264

[217 rows x 4 columns]

```
In [7]: regr_m = linear_model.LinearRegression()
train_x = np.asanyarray(train[['ENGINE_SIZE', 'FUELCONSUMPTION_COMB']])
train_y = np.asanyarray(train[['CO2EMISSIONS']])
regr_m.fit(train_x, train_y)
# The coefficients
print('Coefficients: ', regr_m.coef_)
print('Intercept: ', regr_m.intercept_)
```

Coefficients: [[19.09277508 10.05170185]]
Intercept: [76.17520539]

```
In [8]: test_x = np.asanyarray(test[['ENGINE_SIZE', 'FUELCONSUMPTION_COMB']])
test_y = np.asanyarray(test[['CO2EMISSIONS']])
test_y_ = regr_m.predict(test_x)
print("Mean absolute error: %.2f" % np.mean(np.absolute(test_y_ - test_y)))
print("Residual sum of squares (MSE): %.2f" % np.mean((test_y_ - test_y) ** 2))
print("R2-score: %.2f" % r2_score(test_y_ , test_y) )
```

Mean absolute error: 17.85
Residual sum of squares (MSE): 623.88
R2-score: 0.83

```
In [9]: regr_m = linear_model.LinearRegression()
train_x = np.asanyarray(train[['ENGINE_SIZE', 'CYLINDERS']])
train_y = np.asanyarray(train[['CO2EMISSIONS']])
regr_m.fit(train_x, train_y)
# The coefficients
print('Coefficients: ', regr_m.coef_)
print('Intercept: ', regr_m.intercept_)
```

Coefficients: [[27.2783142 10.24073354]]
Intercept: [105.60906455]

```
In [10]: test_x = np.asanyarray(test[['ENGINE_SIZE', 'CYLINDERS']])
test_y = np.asanyarray(test[['CO2EMISSIONS']])
test_y_ = regr_m.predict(test_x)
print("Mean absolute error: %.2f" % np.mean(np.absolute(test_y_ - test_y)))
print("Residual sum of squares (MSE): %.2f" % np.mean((test_y_ - test_y) ** 2))
print("R2-score: %.2f" % r2_score(test_y_ , test_y) )
```

Mean absolute error: 23.38
Residual sum of squares (MSE): 930.63
R2-score: 0.73

```
In [11]: regr_m = linear_model.LinearRegression()
train_x = np.asanyarray(train[['CYLINDERS', 'FUELCONSUMPTION_COMB']])
train_y = np.asanyarray(train[['CO2EMISSIONS']])
regr_m.fit(train_x, train_y)
# The coefficients
print('Coefficients: ', regr_m.coef_)
print('Intercept: ', regr_m.intercept_)
```

Coefficients: [[14.01904391 10.81328844]]
Intercept: [50.06131774]

```
In [12]: test_x = np.asanyarray(test[['CYLINDERS', 'FUELCONSUMPTION_COMB']])
test_y = np.asanyarray(test[['CO2EMISSIONS']])
test_y_ = regr_m.predict(test_x)
print("Mean absolute error: %.2f" % np.mean(np.absolute(test_y_ - test_y)))
print("Residual sum of squares (MSE): %.2f" % np.mean((test_y_ - test_y) ** 2))
print("R2-score: %.2f" % r2_score(test_y_ , test_y) )
```

Mean absolute error: 18.74
Residual sum of squares (MSE): 659.46
R2-score: 0.83

```
In [13]: regr_m = linear_model.LinearRegression()  
train_x = np.asanyarray(train[['ENGINE_SIZE', 'CYLINDERS', 'FUELCONSUMPTION_COMB']])  
train_y = np.asanyarray(train[['CO2EMISSIONS']])  
regr_m.fit (train_x, train_y)  
# The coefficients  
print ('Coefficients: ', regr_m.coef_)  
print ('Intercept: ',regr_m.intercept_)
```

Coefficients: [[9.77393562 8.20703436 9.87735185]]
Intercept: [61.84041662]

```
In [14]: test_x = np.asanyarray(test[['ENGINE_SIZE', 'CYLINDERS', 'FUELCONSUMPTION_COMB']])  
test_y = np.asanyarray(test[['CO2EMISSIONS']])  
test_y_ = regr_m.predict(test_x)  
print("Mean absolute error: %.2f" % np.mean(np.absolute(test_y_ - test_y)))  
print("Residual sum of squares (MSE): %.2f" % np.mean((test_y_ - test_y) ** 2))  
print("R2-score: %.2f" % r2_score(test_y_ , test_y) )
```

Mean absolute error: 17.97
Residual sum of squares (MSE): 616.56
R2-score: 0.84

In []:

Q 3. Use MLR to estimate the Mileage per gallon (MPG) using Auto-MPG dataset.

```
In [17]: import pandas as pd
import numpy as np
import pylab as pl
import matplotlib.pyplot as plt
from sklearn.metrics import r2_score
from sklearn import linear_model
```

```
In [18]: df = pd.read_csv("auto-mpg.csv")
```

```
In [19]: df.head()
```

```
Out[19]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin	car name
0	18.0	8	307.0	130	3504	12.0	70	1	chevrolet chevelle malibu
1	15.0	8	350.0	165	3693	11.5	70	1	buick skylark 320
2	18.0	8	318.0	150	3436	11.0	70	1	plymouth satellite
3	16.0	8	304.0	150	3433	12.0	70	1	amc rebel sst
4	17.0	8	302.0	140	3449	10.5	70	1	ford torino

```
In [20]: column_name = 'car name'
df = df.drop(column_name, axis=1)
```

```
In [21]: df.head()
```

```
Out[21]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin
0	18.0	8	307.0	130	3504	12.0	70	1
1	15.0	8	350.0	165	3693	11.5	70	1
2	18.0	8	318.0	150	3436	11.0	70	1
3	16.0	8	304.0	150	3433	12.0	70	1
4	17.0	8	302.0	140	3449	10.5	70	1

```
In [22]: cdf = df[['cylinders','displacement','horsepower','weight','acceleration',
                  'model year','origin','mpg']]
```


```
In [23]: cdf.head()
```

```
Out[23]:
```

	cylinders	displacement	horsepower	weight	acceleration	model year	origin	mpg
0	8	307.0	130	3504	12.0	70	1	18.0
1	8	350.0	165	3693	11.5	70	1	15.0
2	8	318.0	150	3436	11.0	70	1	18.0
3	8	304.0	150	3433	12.0	70	1	16.0

	cylinders	displacement	horsepower	weight	acceleration	model year	origin	mpg
4	8	302.0	140	3449	10.5	70	1	17.0

```
In [24]: viz = cdf[['cylinders','displacement','horsepower','weight','acceleration',
                  'model year','origin','mpg']]
viz.hist()
plt.show()
```



```
In [25]: msk = np.random.rand(len(df)) < 0.8
train = cdf[msk]
test = cdf[~msk]
print(train)
print(test)
```

	cylinders	displacement	horsepower	weight	acceleration	model year	\
0	8	307.0	130	3504	12.0	70	
1	8	350.0	165	3693	11.5	70	
2	8	318.0	150	3436	11.0	70	
3	8	304.0	150	3433	12.0	70	
4	8	302.0	140	3449	10.5	70	
..	
393	4	140.0	86	2790	15.6	82	
394	4	97.0	52	2130	24.6	82	
395	4	135.0	84	2295	11.6	82	
396	4	120.0	79	2625	18.6	82	
397	4	119.0	82	2720	19.4	82	

	origin	mpg
0	1	18.0
1	1	15.0
2	1	18.0
3	1	16.0
4	1	17.0
..
393	1	27.0
394	2	44.0
395	1	32.0
396	1	28.0
397	1	31.0

[326 rows x 8 columns]

	cylinders	displacement	horsepower	weight	acceleration	model year	\
10	8	383.0	170	3563	10.0	70	
12	8	400.0	150	3761	9.5	70	
20	4	110.0	87	2672	17.5	70	
22	4	104.0	95	2375	17.5	70	
29	4	97.0	88	2130	14.5	71	
..	


381	4	107.0	75	2205	14.5	82
382	4	108.0	70	2245	16.9	82
384	4	91.0	67	1965	15.7	82
386	6	181.0	110	2945	16.4	82
390	4	144.0	96	2665	13.9	82

	origin	mpg
10	1	15.0
12	1	15.0
20	2	25.0
22	2	25.0
29	3	27.0
..
381	3	36.0
382	3	34.0
384	3	32.0
386	1	25.0
390	3	32.0

[72 rows x 8 columns]

```
In [26]: # Calculate the correlation matrix
import seaborn as sns
correlation_matrix = cdf.corr()

# Create a heatmap using seaborn
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True)
plt.title('Correlation Heatmap')
plt.show()
```



Simple linear regression

```
In [27]: regr = linear_model.LinearRegression()
train_x = np.asanyarray(train[['cylinders']])
train_y = np.asanyarray(train[['mpg']])
```

```

regr.fit (train_x, train_y)
# The coefficients
print ('Coefficients: ', regr.coef_)
print ('Intercept: ',regr.intercept_)

```

Coefficients: $\begin{bmatrix} -3.56892871 \end{bmatrix}$
Intercept: 43.12207679

```

In [28]: test_x = np.asanyarray(test[['cylinders']])
test_y = np.asanyarray(test[['mpg']])
#print(test_y)
test_y_ = regr.predict(test_x)
#print(test_y_)
print("Mean absolute error: %.2f" % np.mean(np.absolute(test_y_ - test_y)))
print("Residual sum of squares (MSE): %.2f" % np.mean((test_y_ - test_y) ** 2))
print("R2-score: %.2f" % r2_score(test_y_ , test_y) )

```

Mean absolute error: 3.15
Residual sum of squares (MSE): 14.93
R2-score: 0.60

```

In [29]: regr = linear_model.LinearRegression()
train_x = np.asanyarray(train[['displacement']])
train_y = np.asanyarray(train[['mpg']])
regr.fit (train_x, train_y)
# The coefficients
print ('Coefficients: ', regr.coef_)
print ('Intercept: ',regr.intercept_)

```

Coefficients: $\begin{bmatrix} -0.06099215 \end{bmatrix}$
Intercept: 35.41953102

```

In [30]: test_x = np.asanyarray(test[['displacement']])
test_y = np.asanyarray(test[['mpg']])
#print(test_y)
test_y_ = regr.predict(test_x)
#print(test_y_)
print("Mean absolute error: %.2f" % np.mean(np.absolute(test_y_ - test_y)))
print("Residual sum of squares (MSE): %.2f" % np.mean((test_y_ - test_y) ** 2))
print("R2-score: %.2f" % r2_score(test_y_ , test_y) )

```

Mean absolute error: 3.27
Residual sum of squares (MSE): 14.90
R2-score: 0.65

```

In [31]: regr = linear_model.LinearRegression()
train_x = np.asanyarray(train[['weight']])
train_y = np.asanyarray(train[['mpg']])
regr.fit (train_x, train_y)
# The coefficients
print ('Coefficients: ', regr.coef_)
print ('Intercept: ',regr.intercept_)

```

Coefficients: $\begin{bmatrix} -0.00770563 \end{bmatrix}$
Intercept: 46.51894909

```

In [32]: test_x = np.asanyarray(test[['weight']])
test_y = np.asanyarray(test[['mpg']])
#print(test_y)
test_y_ = regr.predict(test_x)
#print(test_y_)
print("Mean absolute error: %.2f" % np.mean(np.absolute(test_y_ - test_y)))
print("Residual sum of squares (MSE): %.2f" % np.mean((test_y_ - test_y) ** 2))
print("R2-score: %.2f" % r2_score(test_y_ , test_y) )

```

Mean absolute error: 2.79
Residual sum of squares (MSE): 12.12

R2-score: 0.71

Multiple linear regression

```
In [33]: regr_m = linear_model.LinearRegression()
train_x = np.asanyarray(train[['cylinders', 'displacement', 'weight']])
train_y = np.asanyarray(train[['mpg']])
regr_m.fit (train_x, train_y)
# The coefficients
print ('Coefficients: ', regr_m.coef_)
print ('Intercept: ',regr_m.intercept_)
```

Coefficients: $\begin{bmatrix} -0.17889553 & -0.01419798 & -0.00577035 \end{bmatrix}$
Intercept: $[44.48715502]$

```
In [34]: test_x = np.asanyarray(test[['cylinders', 'displacement', 'weight']])
test_y = np.asanyarray(test[['mpg']])
test_y_ = regr_m.predict(test_x)
print("Mean absolute error: %.2f" % np.mean(np.absolute(test_y_ - test_y)))
print("Residual sum of squares (MSE): %.2f" % np.mean((test_y_ - test_y) ** 2))
print("R2-score: %.2f" % r2_score(test_y_ , test_y) )
```

Mean absolute error: 2.81
Residual sum of squares (MSE): 11.80
R2-score: 0.72

```
In [35]: from sklearn import linear_model
regr_m = linear_model.LinearRegression()
train_x = np.asanyarray(train[['cylinders', 'displacement']])
train_y = np.asanyarray(train[['mpg']])
regr_m.fit (train_x, train_y)
# The coefficients
print ('Coefficients: ', regr_m.coef_)
print ('Intercept: ',regr_m.intercept_)
```

Coefficients: $\begin{bmatrix} -0.33114017 & -0.05586234 \end{bmatrix}$
Intercept: $[36.23523441]$

```
In [36]: test_x = np.asanyarray(test[['cylinders', 'displacement']])
test_y = np.asanyarray(test[['mpg']])
test_y_ = regr_m.predict(test_x)
print("Mean absolute error: %.2f" % np.mean(np.absolute(test_y_ - test_y)))
print("Residual sum of squares (MSE): %.2f" % np.mean((test_y_ - test_y) ** 2))
print("R2-score: %.2f" % r2_score(test_y_ , test_y) )
```

Mean absolute error: 3.25
Residual sum of squares (MSE): 14.69
R2-score: 0.65

```
In [37]: from sklearn import linear_model
regr_m = linear_model.LinearRegression()
train_x = np.asanyarray(train[['displacement', 'weight']])
train_y = np.asanyarray(train[['mpg']])
regr_m.fit (train_x, train_y)
# The coefficients
print ('Coefficients: ', regr_m.coef_)
print ('Intercept: ',regr_m.intercept_)
```

Coefficients: $\begin{bmatrix} -0.0168571 & -0.00578506 \end{bmatrix}$
Intercept: $[44.06846878]$

```
In [38]: test_x = np.asanyarray(test[['displacement', 'weight']])
test_y = np.asanyarray(test[['mpg']])
test_y_ = regr_m.predict(test_x)
print("Mean absolute error: %.2f" % np.mean(np.absolute(test_y_ - test_y)))
print("Residual sum of squares (MSE): %.2f" % np.mean((test_y_ - test_y) ** 2))
print("R2-score: %.2f" % r2_score(test_y_ , test_y) )
```

```
Mean absolute error: 2.81  
Residual sum of squares (MSE): 11.84  
R2-score: 0.72
```

In []:

Comparison tables for FuelConsumption.csv:

Linear regression to predict CO2EMISSIONS:

Attributes	MAE	MSE	R2-Score
ENGINE SIZE	24.39	1050.69	0.64
CYLINDERS	26.04	1112.77	0.63

Multiple linear regression (MLR) to predict CO2EMISSIONS:

Attributes	MAE	MSE	R2-Score
ENGINE SIZE, FUEL CONSUMPTION COMB	17.85	623.88	0.83
ENGINE SIZE, CYLINDERS	23.38	930.63	0.73
CYLINDERS, FUEL CONSUMPTION COMB	18.74	659.46	0.83
ENGINE SIZE, CYLINDERS, FUEL CONSUMPTION COMB	17.97	616.56	0.84

Comparison tables for auto-mpg.csv:

Linear regression to predict mpg:

Attributes	MAE	MSE	R2-Score
cylinders	3.15	14.93	0.60
displacement	3.27	14.90	0.65
weight	2.79	12.12	0.71

Multiple linear regression (MLR) to predict mpg:

Attributes	MAE	MSE	R2-Score
cylinders, displacement, weight	2.81	11.80	0.72
cylinders, displacement	3.25	14.69	0.65
displacement, weight	2.81	11.84	0.72