

Name - Arkadipta Mojumder
Registration Number - 22MCA0201
Digital Assessment 2

13-07-2023-penguin-classification

July 26, 2023

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report
from sklearn.tree import DecisionTreeClassifier
```

```
[ ]: data = pd.read_csv('penguins_size.csv')
```


```
[ ]: data.head()
```

```
[ ]: species      island  culmen_length_mm  culmen_depth_mm  flipper_length_mm  \
0  Adelie  Torgersen           39.1           18.7           181.0
1  Adelie  Torgersen           39.5           17.4           186.0
2  Adelie  Torgersen           40.3           18.0           195.0
3  Adelie  Torgersen           NaN           NaN           NaN
4  Adelie  Torgersen           36.7           19.3           193.0

      body_mass_g      sex
0          3750.0    MALE
1          3800.0  FEMALE
2          3250.0  FEMALE
3              NaN      NaN
4          3450.0  FEMALE
```


```
[ ]: data.species.value_counts().plot(kind='pie')
```

```
[ ]: <Axes: ylabel='species'>
```



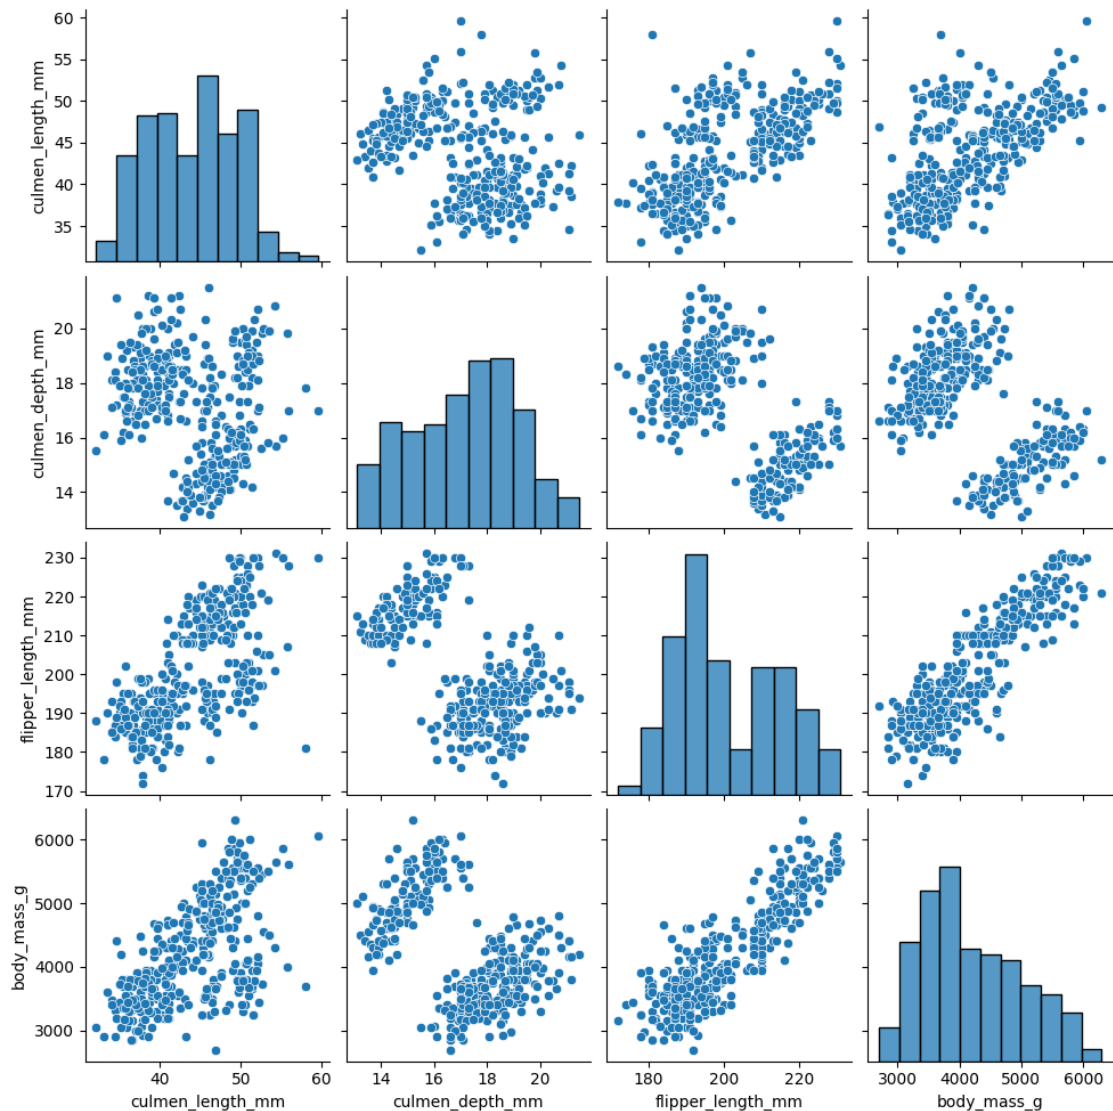
```
[ ]: import seaborn as sns
sns.countplot(data=data,x='species',hue='sex')
```

```
[ ]: <Axes: xlabel='species', ylabel='count'>
```



```
[ ]: sns.pairplot(data)
```

```
[ ]: <seaborn.axisgrid.PairGrid at 0x7a0e4c09c490>
```



```
[ ]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 344 entries, 0 to 343
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   species               344 non-null   object
1   island                344 non-null   object
2   culmen_length_mm      342 non-null   float64
3   culmen_depth_mm       342 non-null   float64
4   flipper_length_mm     342 non-null   float64
5   body_mass_g           342 non-null   float64
```

```
6    sex          334 non-null    object
dtypes: float64(4), object(3)
memory usage: 18.9+ KB
```

```
[ ]: data.describe()
```

```
[ ]:      culmen_length_mm  culmen_depth_mm  flipper_length_mm  body_mass_g
count      342.000000      342.000000      342.000000      342.000000
mean        43.921930       17.151170       200.915205     4201.754386
std         5.459584        1.974793        14.061714      801.954536
min         32.100000       13.100000       172.000000     2700.000000
25%         39.225000       15.600000       190.000000     3550.000000
50%         44.450000       17.300000       197.000000     4050.000000
75%         48.500000       18.700000       213.000000     4750.000000
max         59.600000       21.500000       231.000000     6300.000000
```

```
[ ]: data.isnull().sum()
```

```
[ ]: species          0
     island           0
     culmen_length_mm  2
     culmen_depth_mm   2
     flipper_length_mm  2
     body_mass_g       2
     sex              10
dtype: int64
```

```
[ ]: data.dropna(subset=['sex'],inplace=True)
     data.isnull().sum()
```

```
[ ]: species          0
     island           0
     culmen_length_mm  0
     culmen_depth_mm   0
     flipper_length_mm  0
     body_mass_g       0
     sex              0
dtype: int64
```

```
[ ]: data['species'].value_counts()
```

```
[ ]: Adelie      146
     Gentoo     120
     Chinstrap   68
     Name: species, dtype: int64
```

```
[ ]: data['sex'].value_counts()
```

```
[ ]: MALE      168
      FEMALE   165
      .         1
      Name: sex, dtype: int64
```

```
[ ]: data=data[data.sex!='.']
```

```
[ ]: data['sex'].value_counts()
```

```
[ ]: MALE      168
      FEMALE   165
      Name: sex, dtype: int64
```

```
[ ]: x=pd.get_dummies(data.drop('species',axis=1),columns=['island','sex'])
```

```
[ ]: x.head()
```

```
[ ]:      culmen_length_mm  culmen_depth_mm  flipper_length_mm  body_mass_g  \
0          39.1           18.7           181.0        3750.0
1          39.5           17.4           186.0        3800.0
2          40.3           18.0           195.0        3250.0
4          36.7           19.3           193.0        3450.0
5          39.3           20.6           190.0        3650.0
```

	island_Biscoe	island_Dream	island_Torgersen	sex_FEMALE	sex_MALE
0	0	0	1	0	1
1	0	0	1	1	0
2	0	0	1	1	0
4	0	0	1	1	0
5	0	0	1	0	1

```
[ ]: y=data['species']
```

```
[ ]: from sklearn.preprocessing import LabelEncoder
```

```
[ ]: le=LabelEncoder()
      y=le.fit_transform(y)
      y
```

[illegible]

```

1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2,
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
2, 2, 2])

```

```
[ ]: from sklearn.model_selection import train_test_split
```

```
[ ]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2)
```

```
[ ]: result=pd.DataFrame()
Models=[]
Accuracy=[]
```

```
[ ]: lr=LogisticRegression()
```

```
[ ]: knn=KNeighborsClassifier()
k_pred=knn.fit(x_train,y_train).predict(x_test)
print("Accuracy={}".format(accuracy_score(y_test,k_pred)))
```

Accuracy=0.8507462686567164

```
[ ]: kvl=range(1,10)
bestk=0
acc=0
b_acc=0
for k in kvl:
    knn=KNeighborsClassifier(n_neighbors=k)
    k_pred=knn.fit(x_train,y_train).predict(x_test)
    acc=accuracy_score(y_test,k_pred)
    if acc>b_acc:
        b_acc=acc
        bestk=k
print("Max accuracy:", b_acc)
print("Best k:", bestk)
```

Max accuracy: 0.8656716417910447

Best k: 1

```
[ ]: knn = KNeighborsClassifier(n_neighbors=bestk)
knn_pred = knn.fit(x_train, y_train).predict(x_test)
Models.append('K-Nearest Neighbor')
Accuracy.append(accuracy_score(y_test, knn_pred))
```

```
print("Accuracy score using K-Nearest Neighbor is: {}".format(accuracy_score(y_test, knn_pred)*100))
print(classification_report(y_test, knn_pred))
```

Accuracy score using K-Nearest Neighbor is: 86.56716417910447%

	precision	recall	f1-score	support
0	0.94	0.81	0.87	37
1	0.64	0.90	0.75	10
2	0.90	0.95	0.93	20
accuracy			0.87	67
macro avg	0.83	0.89	0.85	67
weighted avg	0.88	0.87	0.87	67

```
[ ]: dst = DecisionTreeClassifier()
dst_pred = dst.fit(x_train, y_train).predict(x_test)
Models.append('Decision Trees')
Accuracy.append(accuracy_score(y_test, dst_pred))
print("Accuracy score using Decision Trees is: {}".format(accuracy_score(y_test, dst_pred)*100))
print(classification_report(y_test, dst_pred))
```

Accuracy score using Decision Trees is: 95.52238805970148%

	precision	recall	f1-score	support
0	0.97	0.95	0.96	37
1	0.91	1.00	0.95	10
2	0.95	0.95	0.95	20
accuracy			0.96	67
macro avg	0.94	0.97	0.95	67
weighted avg	0.96	0.96	0.96	67

```
[ ]: result['models']=Models
result['accuracy']=Accuracy
```

```
[ ]: result
```

```
[ ]:
      models accuracy
0  K-Nearest Neighbor 0.865672
1    Decision Trees 0.955224
```


breast cancer-diagnosis-using-ml techniques

July 26, 2023

```
[ ]: import numpy as np
import pandas as pd
from pandas import Series, DataFrame
import seaborn as sns
import matplotlib.pyplot as plt
palette = 'magma'

%matplotlib inline
```

```
[ ]: data = pd.read_csv("data.csv")
data.head()
```

```
[ ]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	842302	M	17.99	10.38	122.80	1001.0	
1	842517	M	20.57	17.77	132.90	1326.0	
2	84300903	M	19.69	21.25	130.00	1203.0	
3	84348301	M	11.42	20.38	77.58	386.1	
4	84358402	M	20.29	14.34	135.10	1297.0	

	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	\
0	0.11840	0.27760	0.3001	0.14710	
1	0.08474	0.07864	0.0869	0.07017	
2	0.10960	0.15990	0.1974	0.12790	
3	0.14250	0.28390	0.2414	0.10520	
4	0.10030	0.13280	0.1980	0.10430	

...	texture_worst	perimeter_worst	area_worst	smoothness_worst	\
0	17.33	184.60	2019.0	0.1622	
1	23.41	158.80	1956.0	0.1238	
2	25.53	152.50	1709.0	0.1444	
3	26.50	98.87	567.7	0.2098	
4	16.67	152.20	1575.0	0.1374	

	compactness_worst	concavity_worst	concave points_worst	symmetry_worst	\
0	0.6656	0.7119	0.2654	0.4601	
1	0.1866	0.2416	0.1860	0.2750	
2	0.4245	0.4504	0.2430	0.3613	

3	0.8663	0.6869	0.2575	0.6638
4	0.2050	0.4000	0.1625	0.2364

	fractal_dimension_worst	Unnamed: 32
0	0.11890	NaN
1	0.08902	NaN
2	0.08758	NaN
3	0.17300	NaN
4	0.07678	NaN

[5 rows x 33 columns]

```
[ ]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     569 non-null    int64
1   diagnosis                             569 non-null    object
2   radius_mean                           569 non-null    float64
3   texture_mean                           569 non-null    float64
4   perimeter_mean                         569 non-null    float64
5   area_mean                             569 non-null    float64
6   smoothness_mean                       569 non-null    float64
7   compactness_mean                      569 non-null    float64
8   concavity_mean                        569 non-null    float64
9   concave points_mean                   569 non-null    float64
10  symmetry_mean                         569 non-null    float64
11  fractal_dimension_mean                569 non-null    float64
12  radius_se                             569 non-null    float64
13  texture_se                             569 non-null    float64
14  perimeter_se                           569 non-null    float64
15  area_se                               569 non-null    float64
16  smoothness_se                         569 non-null    float64
17  compactness_se                        569 non-null    float64
18  concavity_se                          569 non-null    float64
19  concave points_se                     569 non-null    float64
20  symmetry_se                           569 non-null    float64
21  fractal_dimension_se                  569 non-null    float64
22  radius_worst                          569 non-null    float64
23  texture_worst                         569 non-null    float64
24  perimeter_worst                       569 non-null    float64
25  area_worst                            569 non-null    float64
26  smoothness_worst                      569 non-null    float64
27  compactness_worst                     569 non-null    float64
```

```

28 concavity_worst      569 non-null    float64
29 concave points_worst  569 non-null    float64
30 symmetry_worst       569 non-null    float64
31 fractal_dimension_worst 569 non-null    float64
32 Unnamed: 32          0 non-null      float64
dtypes: float64(31), int64(1), object(1)
memory usage: 146.8+ KB

```

Removing non-essential column

```
[ ]: data.drop("id", axis = 1, inplace = True)
```

```
[ ]: data.head()
```

```
[ ]:
diagnosis  radius_mean  texture_mean  perimeter_mean  area_mean  \
0         M         17.99         10.38         122.80        1001.0
1         M         20.57         17.77         132.90        1326.0
2         M         19.69         21.25         130.00        1203.0
3         M         11.42         20.38          77.58         386.1
4         M         20.29         14.34         135.10        1297.0

smoothness_mean  compactness_mean  concavity_mean  concave points_mean  \
0         0.11840         0.27760         0.3001         0.14710
1         0.08474         0.07864         0.0869         0.07017
2         0.10960         0.15990         0.1974         0.12790
3         0.14250         0.28390         0.2414         0.10520
4         0.10030         0.13280         0.1980         0.10430

symmetry_mean  ... texture_worst  perimeter_worst  area_worst  \
0         0.2419  ...         17.33         184.60        2019.0
1         0.1812  ...         23.41         158.80        1956.0
2         0.2069  ...         25.53         152.50        1709.0
3         0.2597  ...         26.50          98.87         567.7
4         0.1809  ...         16.67         152.20        1575.0

smoothness_worst  compactness_worst  concavity_worst  concave points_worst  \
0         0.1622         0.6656         0.7119         0.2654
1         0.1238         0.1866         0.2416         0.1860
2         0.1444         0.4245         0.4504         0.2430
3         0.2098         0.8663         0.6869         0.2575
4         0.1374         0.2050         0.4000         0.1625

symmetry_worst  fractal_dimension_worst  Unnamed: 32
0         0.4601         0.11890         NaN
1         0.2750         0.08902         NaN
2         0.3613         0.08758         NaN
3         0.6638         0.17300         NaN
4         0.2364         0.07678         NaN

```

[5 rows x 32 columns]

```
[ ]: data = data[['diagnosis', 'radius_mean', 'texture_mean', 'perimeter_mean',  
                'area_mean', 'smoothness_mean', 'compactness_mean', 'concavity_mean',  
                'concave points_mean']]
```

Converting Diagnosis to Dummy Variables

```
[ ]: malignant = pd.get_dummies(data['diagnosis'], drop_first=True)  
malignant
```

```
[ ]:      M  
0      1  
1      1  
2      1  
3      1  
4      1  
... ..  
564    1  
565    1  
566    1  
567    1  
568    0
```

[569 rows x 1 columns]

```
[ ]: data = pd.concat([data, malignant], axis=1)
```

```
[ ]: data.drop('diagnosis', axis=1, inplace=True)  
data.head()
```

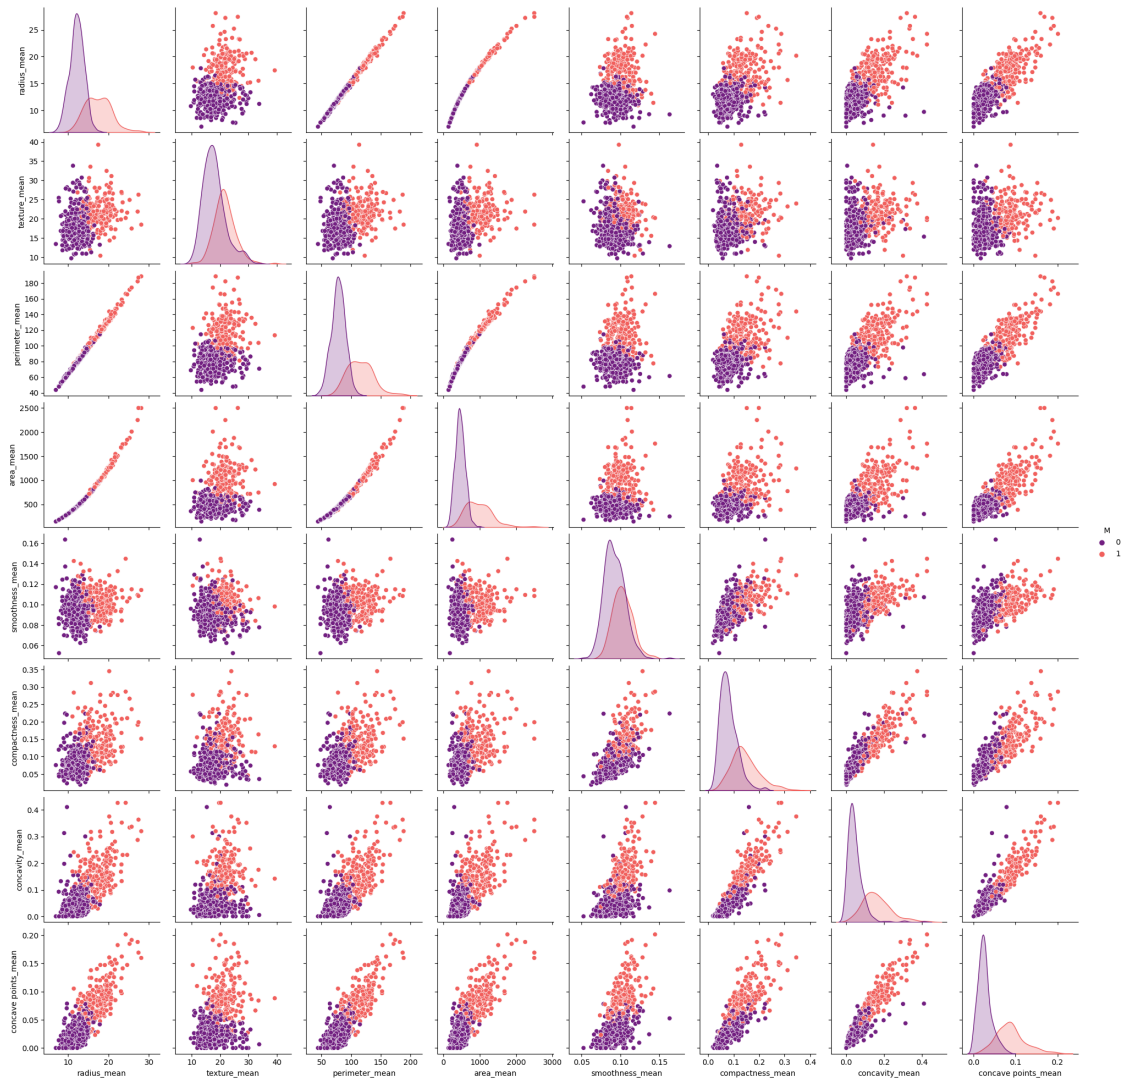
```
[ ]:      radius_mean  texture_mean  perimeter_mean  area_mean  smoothness_mean  \  
0          17.99         10.38         122.80       1001.0         0.11840  
1          20.57         17.77         132.90       1326.0         0.08474  
2          19.69         21.25         130.00       1203.0         0.10960  
3          11.42         20.38          77.58        386.1         0.14250  
4          20.29         14.34         135.10       1297.0         0.10030  
  
      compactness_mean  concavity_mean  concave points_mean  M  
0          0.27760         0.3001         0.14710  1  
1          0.07864         0.0869         0.07017  1  
2          0.15990         0.1974         0.12790  1  
3          0.28390         0.2414         0.10520  1  
4          0.13280         0.1980         0.10430  1
```

Exploratory Data Analysis

```
[ ]: plt.figure(figsize=(10,6))
sns.pairplot(data, hue='M', palette=palette)
```


```
[ ]: <seaborn.axisgrid.PairGrid at 0x7ab78c10c820>
```

<Figure size 1000x600 with 0 Axes>



```
[ ]: sns.heatmap(data.corr(), annot=True, cmap=palette)
```

```
[ ]: <Axes: >
```



Standardize the variables

```
[ ]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(data.drop('M', axis=1))
```

```
[ ]: StandardScaler()
```

```
[ ]: scaled_features = scaler.transform(data.drop('M', axis=1))
```

```
[ ]: df_feat = pd.DataFrame(scaled_features, columns=data.columns[:-1])
df_feat.head()
```

```
[ ]:      radius_mean  texture_mean  perimeter_mean  area_mean  smoothness_mean  \
0      1.097064    -2.073335      1.269934    0.984375      1.568466
1      1.829821    -0.353632      1.685955    1.908708     -0.826962
2      1.579888     0.456187      1.566503    1.558884     0.942210
3     -0.768909     0.253732     -0.592687   -0.764464     3.283553
```

4	1.750297	-1.151816	1.776573	1.826229	0.280372
---	----------	-----------	----------	----------	----------

	compactness_mean	concavity_mean	concave	points_mean
0	3.283515	2.652874		2.532475
1	-0.487072	-0.023846		0.548144
2	1.052926	1.363478		2.037231
3	3.402909	1.915897		1.451707
4	0.539340	1.371011		1.428493

Split the Data into Training and Testing Set

```
[ ]: from sklearn.model_selection import train_test_split
```

```
[ ]: X = df_feat
y = data['M']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↪random_state=101)
```

Prediction and Evaluation

```
[ ]: from sklearn.neighbors import KNeighborsClassifier
```

Create a loop to train KNN models with different k values and track the error rate for each model in a list.


```
[ ]: error_rate = []
for i in range(1,40):

    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(X_train, y_train)
    pred_i = knn.predict(X_test)
    error_rate.append(np.mean(pred_i != y_test))
```

Generate a plot based on the data in the list

```
[ ]: plt.figure(figsize=(10,6))
plt.plot(range(1,40), error_rate, color='blue', linestyle='dashed', marker='o',
        markerfacecolor='red', markersize=10)
plt.title('Error Rate VS K Value')
plt.xlabel('K Value')
plt.ylabel('Error Rate')
```

```
[ ]: Text(0, 0.5, 'Error Rate')
```



0.0.1 Set k to 17, then train the model and make predictions.

```
[ ]: knn = KNeighborsClassifier(n_neighbors=17)
      knn.fit(X_train, y_train)
      pred = knn.predict(X_test)
```

Evaluate the model using a classification report and a confusion matrix

```
[ ]: from sklearn.metrics import classification_report, confusion_matrix
```

```
[ ]: print(confusion_matrix(y_test, pred))
```

```
[[100  5]
 [ 6 60]]
```

```
[ ]: print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
0	0.94	0.95	0.95	105
1	0.92	0.91	0.92	66
accuracy			0.94	171
macro avg	0.93	0.93	0.93	171

weighted avg	0.94	0.94	0.94	171
--------------	------	------	------	-----