

Name - Arkadipta Mojumder

Registration Number - 22MCA0201

Task 4 - CNN on Penguin dataset

```
In [ ]: import pandas as pd
import numpy as np
import os
# Plotting libraries
import matplotlib.pyplot as plt
# SKLearn libraries
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
# Tensorflow libraries
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
```

```
In [ ]: penguin_data = pd.read_csv("penguins.csv")
```

Preprocessing

```
In [ ]: penguin_data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 344 entries, 0 to 343
Data columns (total 9 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   Unnamed: 0           344 non-null   int64
 1   species              344 non-null   object
 2   island               344 non-null   object
 3   bill_length_mm       342 non-null   float64
 4   bill_depth_mm        342 non-null   float64
 5   flipper_length_mm    342 non-null   float64
 6   body_mass_g          342 non-null   float64
 7   sex                  333 non-null   object
 8   year                 344 non-null   int64
dtypes: float64(4), int64(2), object(3)
memory usage: 24.3+ KB
```

```
In [ ]: penguin_data.head(10)
```

```
Out[ ]:      Unnamed: 0  species  island  bill_length_mm  bill_depth_mm  flipper_length_mm  body_mass_g  sex  year
0              0      Adelle  Torgersen         39.1           18.7           181.0         3750.0  male  2007
1              1              2      Adelle  Torgersen         39.5           17.4           186.0         3800.0  female 2007
2              2              3      Adelle  Torgersen         40.3           18.0           195.0         3250.0  female 2007
3              3              4      Adelle  Torgersen         NaN             NaN             NaN             NaN      NaN  2007
4              4              5      Adelle  Torgersen         36.7           19.3           193.0         3450.0  female 2007
5              5              6      Adelle  Torgersen         39.3           20.6           190.0         3650.0  male  2007
6              6              7      Adelle  Torgersen         38.9           17.8           181.0         3625.0  female 2007
7              7              8      Adelle  Torgersen         39.2           19.6           195.0         4675.0  male  2007
8              8              9      Adelle  Torgersen         34.1           18.1           193.0         3475.0  NaN  2007
9              9             10      Adelle  Torgersen         42.0           20.2           190.0         4250.0  NaN  2007
```

```
In [ ]: penguin_data = penguin_data.dropna()
penguin_data = penguin_data[penguin_data['sex'] != '.']
penguin_data.isna().any()
```

```
Out[ ]: Unnamed: 0      False
species           False
island            False
bill_length_mm    False
bill_depth_mm     False
flipper_length_mm False
body_mass_g       False
sex               False
year              False
dtype: bool
```

Preparing dataset

```
In [ ]: #One hot encoding
one_hot_encoded = pd.get_dummies(penguin_data[['sex', 'island']])
penguin_data.drop(['island', 'sex'], axis = 1, inplace = True)
penguin_data = pd.concat([penguin_data, one_hot_encoded], axis = 1)
penguin_data
```

```
Out[ ]:      Unnamed: 0  species  bill_length_mm  bill_depth_mm  flipper_length_mm  body_mass_g  year  sex_female  sex_male  islan
0              0      Adelle         39.1           18.7           181.0         3750.0  2007           0           1
1              1              2      Adelle         39.5           17.4           186.0         3800.0  2007           1           0
2              2              3      Adelle         40.3           18.0           195.0         3250.0  2007           1           0
3              3              4      Adelle         36.7           19.3           193.0         3450.0  2007           1           0
4              4              5      Adelle         39.3           20.6           190.0         3650.0  2007           0           1
...          ...          ...          ...          ...          ...          ...          ...          ...          ...
339          340  Chinstrap         55.8           19.8           207.0         4000.0  2009           0           1
340          341  Chinstrap         43.5           18.1           202.0         3400.0  2009           1           0
341          342  Chinstrap         49.6           18.2           193.0         3775.0  2009           0           1
342          343  Chinstrap         50.8           19.0           210.0         4100.0  2009           0           1
343          344  Chinstrap         50.2           18.7           198.0         3775.0  2009           1           0
```

333 rows x 12 columns

```
In [ ]: X = penguin_data.loc[:, penguin_data.columns != 'species']
y = penguin_data.loc[:, ['species']]
```

```
In [ ]: y_end = LabelEncoder().fit_transform(y)
y_label = tf.keras.utils.to_categorical(y_end)

/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.py:116: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
```

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y_label, test_size=0.3)
```

```
In [ ]: print(f"Train shape : {X_train.shape}, Y Train : {y_train.shape}")
print(X_train.shape[1:])

Train shape : (233, 11), Y Train : (233, 3)
(11,)
```

```
In [ ]: def get_model():
    model = Sequential([
        keras.layers.Input(shape=X_train.shape[1:]),
        keras.layers.Dense(500, activation='relu'),
        keras.layers.Dense(200, activation='relu'),
        keras.layers.Dense(3, activation='softmax')
    ])

    return model
```

```
In [ ]: model = get_model()

# Compile the model
model.compile(optimizer='adam',
              loss=keras.losses.CategoricalCrossentropy(),
              metrics=['accuracy'])
```

```
In [ ]: model.summary()

Model: "sequential"

Layer (type)                Output Shape                Param #
=====
dense (Dense)                (None, 500)                 6000
dense_1 (Dense)              (None, 200)                100200
dense_2 (Dense)              (None, 3)                  603
=====
Total params: 106,803
Trainable params: 106,803
Non-trainable params: 0
```

```
In [ ]: history = model.fit(X_train, y_train, epochs=30, validation_data=(X_test, y_test), verbose = 1)

Epoch 1/30
8/8 [=====] - 2s 47ms/step - loss: 242.2726 - accuracy: 0.3519 - val_loss: 198.9735 - val_accuracy: 0.3200
Epoch 2/30
8/8 [=====] - 0s 12ms/step - loss: 116.9643 - accuracy: 0.5837 - val_loss: 63.3924 - val_accuracy: 0.2700
Epoch 3/30
8/8 [=====] - 0s 13ms/step - loss: 44.6350 - accuracy: 0.4936 - val_loss: 16.9714 - val_accuracy: 0.6800
Epoch 4/30
8/8 [=====] - 0s 11ms/step - loss: 10.0658 - accuracy: 0.6738 - val_loss: 10.8250 - val_accuracy: 0.7400
Epoch 5/30
8/8 [=====] - 0s 17ms/step - loss: 6.4248 - accuracy: 0.7983 - val_loss: 0.6606 - val_accuracy: 0.9400
Epoch 6/30
8/8 [=====] - 0s 24ms/step - loss: 3.2928 - accuracy: 0.8670 - val_loss: 2.1661 - val_accuracy: 0.9400
Epoch 7/30
8/8 [=====] - 0s 18ms/step - loss: 3.0175 - accuracy: 0.8670 - val_loss: 2.7187 - val_accuracy: 0.8000
Epoch 8/30
8/8 [=====] - 0s 16ms/step - loss: 2.5720 - accuracy: 0.8755 - val_loss: 2.4593 - val_accuracy: 0.9300
Epoch 9/30
8/8 [=====] - 0s 17ms/step - loss: 2.1904 - accuracy: 0.9313 - val_loss: 2.0009 - val_accuracy: 0.9400
Epoch 10/30
8/8 [=====] - 0s 16ms/step - loss: 2.2948 - accuracy: 0.9142 - val_loss: 0.9171 - val_accuracy: 0.9600
Epoch 11/30
8/8 [=====] - 0s 18ms/step - loss: 2.2828 - accuracy: 0.9185 - val_loss: 0.4686 - val_accuracy: 0.9600
Epoch 12/30
8/8 [=====] - 0s 17ms/step - loss: 1.8696 - accuracy: 0.9356 - val_loss: 1.3387 - val_accuracy: 0.9600
Epoch 13/30
8/8 [=====] - 0s 16ms/step - loss: 1.6138 - accuracy: 0.9399 - val_loss: 0.3163 - val_accuracy: 0.9600
Epoch 14/30
8/8 [=====] - 0s 17ms/step - loss: 1.0763 - accuracy: 0.9485 - val_loss: 3.1587 - val_accuracy: 0.9000
Epoch 15/30
8/8 [=====] - 0s 17ms/step - loss: 2.5921 - accuracy: 0.9099 - val_loss: 2.9179 - val_accuracy: 0.9600
Epoch 16/30
8/8 [=====] - 0s 17ms/step - loss: 1.8278 - accuracy: 0.9227 - val_loss: 0.3985 - val_accuracy: 0.9700
Epoch 17/30
8/8 [=====] - 0s 16ms/step - loss: 1.6874 - accuracy: 0.9356 - val_loss: 1.1916 - val_accuracy: 0.9000
Epoch 18/30
8/8 [=====] - 0s 17ms/step - loss: 3.3934 - accuracy: 0.9013 - val_loss: 0.5705 - val_accuracy: 0.9600
Epoch 19/30
8/8 [=====] - 0s 18ms/step - loss: 1.9237 - accuracy: 0.8927 - val_loss: 0.2916 - val_accuracy: 0.9700
Epoch 20/30
8/8 [=====] - 0s 17ms/step - loss: 0.8929 - accuracy: 0.9270 - val_loss: 0.2008 - val_accuracy: 0.9800
Epoch 21/30
8/8 [=====] - 0s 17ms/step - loss: 1.3977 - accuracy: 0.9185 - val_loss: 0.3246 - val_accuracy: 0.9700
Epoch 22/30
8/8 [=====] - 0s 18ms/step - loss: 0.6355 - accuracy: 0.9485 - val_loss: 0.9737 - val_accuracy: 0.9500
Epoch 23/30
8/8 [=====] - 0s 18ms/step - loss: 5.1764 - accuracy: 0.8326 - val_loss: 6.6045 - val_accuracy: 0.7600
Epoch 24/30
8/8 [=====] - 0s 15ms/step - loss: 7.7603 - accuracy: 0.7554 - val_loss: 3.3145 - val_accuracy: 0.8700
Epoch 25/30
8/8 [=====] - 0s 19ms/step - loss: 2.7231 - accuracy: 0.9313 - val_loss: 1.8126 - val_accuracy: 0.9300
Epoch 26/30
8/8 [=====] - 0s 22ms/step - loss: 1.1055 - accuracy: 0.9485 - val_loss: 1.0470 - val_accuracy: 0.9400
Epoch 27/30
8/8 [=====] - 0s 14ms/step - loss: 0.8515 - accuracy: 0.9528 - val_loss: 0.4209 - val_accuracy: 0.9700
Epoch 28/30
8/8 [=====] - 0s 14ms/step - loss: 0.6958 - accuracy: 0.9657 - val_loss: 0.1836 - val_accuracy: 0.9800
Epoch 29/30
8/8 [=====] - 0s 13ms/step - loss: 1.0503 - accuracy: 0.9313 - val_loss: 0.9005 - val_accuracy: 0.9500
Epoch 30/30
8/8 [=====] - 0s 13ms/step - loss: 0.7413 - accuracy: 0.9442 - val_loss: 0.9861 - val_accuracy: 0.9500
```

```
In [ ]: model.evaluate(X_test, y_test)
```

```
4/4 [=====] - 0s 4ms/step - loss: 0.9861 - accuracy: 0.9500
```

```
Out[ ]: [0.9861079454421997, 0.94999998079071]
```

Performance Monitor

```
In [ ]: pd.DataFrame(history.history).plot(figsize = (10,6))
plt.grid(True)
plt.gca().set_ylim(0, 1)
plt.show()
```

